

Preventing Signedness Errors in Numerical Computations in Java

by

Christopher A. Mackie

Supervised by **Michael D. Ernst**

A senior thesis submitted in partial fulfillment of
the requirements for the degree of

**Bachelor of Science
With Departmental Honors
Computer Science & Engineering
University of Washington**

June 2017

Presentation of work given on _____

Thesis and presentation approved by _____

Date _____

ABSTRACT

We have developed and implemented a type system, the Signedness Type System, that captures usage of signed and unsigned integers in Java programs. This type system enables developers to detect errors regarding unsigned integers at compile time, and guarantees that such errors cannot occur at run time. An implementation of this type system for Java, the Signedness Checker, uncovered previously unknown bugs in a case study. The Signedness Checker is available publicly bundled with the Checker Framework (<http://CheckerFramework.org/>).

CONTENTS

Contents	2
1Problem and Motivation	3
2Approach and Uniqueness	4
3Precision Improvements	7
3.1 Masked Shifts	7
3.2 Casted Shifts	7
3.3 Shift Propagation	7
3.4 Fully Compatible Integers	8
4Background and Implementation	9
5Results and Contributions	10
5.1 First Approach	10
5.2 Second Approach	10
5.3 Tool Evaluation	12
6Related Work	14
7Future Work	14
8Conclusions	14
References	14

1 PROBLEM AND MOTIVATION

In order to represent integers, computers make use of several different binary encodings. The two most prevalent of these encodings are unsigned, and two’s complement. The former is used to represent strictly nonnegative integers whereas the latter is used to represent signed integers. While the truth about two’s complement is slightly more complex, for our purposes we can distinguish between unsigned and signed integers by considering how the most significant bit is interpreted. Signed integers use the most significant bit to represent the sign of the integer, whereas unsigned integers utilize all bits to represent magnitude. As a result, unsigned integers can cover a larger range of nonnegative integer values than signed integers given the same number of bits, but they are incapable of representing negative integers. While signed integers are much more common, unsigned integers are more efficient and often useful in low level code, so many programming languages support them; for example, Java 8 introduced utility methods for unsigned integers [3]. However, they are also error-prone: using one where a signed integer is expected, or performing certain arithmetic operations on unsigned integers, can lead to unexpected results.

For the duration of this paper, we call an operator “insensitive” if it produces a correct signed result when run on two signed values and a correct unsigned result when run on two unsigned values. We call an operator “sensitive” if it must be implemented differently for signed and unsigned operands. That is a sensitive operator implemented for signed operands produces incorrect results when run on two unsigned values, and vice versa. For such operations, a programmer must run a different implementation depending on whether the operands are signed or unsigned. Specifically in Java, all sensitive operators are naturally implemented for signed integers, so programmers must make use of utility functions to operate on unsigned integers. See Figure 1 for an example of insensitive and sensitive operators. In this example we learn that subtraction is insensitive, but division is sensitive. In Java, the division operator is implemented for signed operands. In order to apply division to unsigned integers, the programmer must call `Integer::divideUnsigned`, a utility introduced in Java 8.

Misuse of unsigned values can be categorized as follows:

- Using a sensitive operator with unsigned operands.
- Mixing signed and unsigned arguments in a computation.

See Figure 2 for an example of misuses of unsigned values. We see that the main issue with mixing signed and unsigned values is it is not clear how we should interpret the result. Languages like C resolve this issue somewhat arbitrarily by promoting signed integers to unsigned integers. However, this can lead to extremely frustrating runtime errors. As a result, we choose to completely disallow the mixing of signed and unsigned integers.

The first line of defense against most bugs is the compiler. When the compiler is unable to catch bugs, it falls on the programmer to identify and eliminate them, which is prone to human error. The Java compiler is not helpful in finding bugs related to using unsigned numbers because Java’s unsigned integers are supported by a library rather than built into the core language. As a result, Java interprets all integers as signed, and does not issue any errors for the misuses discussed above. We aim to develop a general analysis for verifying the correct usage of unsigned integers. We will then implement that analysis for Java in the form of a compile time signedness checker.

```

\\ Signed: -3, Unsigned: 253
byte a = 0xFFFD;

\\ Signed: 2, Unsigned: 2
byte b = 0x0002;

\\ - is an insensitive operator and thus either
\\ interpretation is consistent
\\ Signed: -1, Unsigned: 255
byte sum = a - b;

\\ / is a sensitive operator and thus only the
\\ operator interpretation is always consistent
\\ (in this case signed)
\\ Signed: -1, Unsigned: 255 (expected 126)
byte div = a / b;

```

Fig. 1. An example of the distinction between insensitive and sensitive operators.

2 APPROACH AND UNIQUENESS

Our approach to detecting and preventing signedness errors is to use a type system. This has a number of benefits. Firstly, compile-time checking permits developers to catch bugs before they become problems for end-users. Secondly, static checking provides a guarantee across all program executions. Finally, type systems are familiar to programmers, who understand how to use them and interpret their warning messages.

We have defined a type system and implemented it for Java in a tool called the Signedness Checker. This type system has five type qualifiers, as seen in Figure 3. Each is written together with a base type; for example, @Unsigned int i declares a variable named i of type @Unsigned int. The specific semantics of each qualifier in the system follows:

- @UnknownSignedness indicates a value for which our tool has no estimate. (This usually leads to a warning.) It is also used for non-numeric values, which our system ignores.
- @Unsigned signifies that a value is unsigned.
- @Signed signifies that a value is signed.
- @Constant is for values known at compile time, such as manifest literals. The programmer might intend to use them as signed or as unsigned values. Even a negative literal may be used as a placeholder for a large positive unsigned integer.
- @SignednessBottom indicates dead code or the null value.

We were initially surprised that it was necessary to include the @Constant qualifier in our system. As it turns out, this is extremely useful for the precision of our type system. It allows users to represent bit patterns however they please, and it allows us to represent values which can reasonably be interpreted as signed or unsigned.

```

\\ Using unsigned integers in a sensitive operator
\\ can lead to incorrect results
\\ Unsigned: 255
byte a = 0xFF;

\\Unsigned: 254
byte b = 0xFE;

\\Unsigned: 0 (expected 1)
byte div = a / b;

\\ Mixing signed and unsigned integers leads to
\\ incomprehensible results
\\ Signed: 1
byte s = 0x01;

\\ Unsigned: 127
byte u = 0x7F;

\\ Signed: -128, Unsigned: 128
byte sum = s + u;

\\ How do we even interpret this?
\\ Do we print -128 or 128?
System.out.println(sum);

```

Fig. 2. An example of possible misuses of unsigned integers.

The programmer writes @Unsigned type qualifiers on type uses in the program's source code, and may sometimes write @Constant on constant fields. Unannotated Java types are given a default qualifier. We can infer the type of unannotated variables by using type introduction rules. The type introduction rules for the Signedness Type System are as follows:

- (1) If the user wrote a type qualifier, use it.
- (2) If an integral expression is a constant at compile time, use @Constant. Integral types are char, byte, short, int, and long.
- (3) Other integral expressions use @Signed.
- (4) Binary integral operators have the least upper bound of the types of its operands as its type.
- (5) Non-integral expressions use @UnknownSignedness.

Some of these type introduction rules require further explanation. As previously stated, we wish to allow users to represent signed and unsigned literals however they please, so we infer all manifest literals as @Constant to allow them to be used as @Signed or @Unsigned. All integral variables will be treated by Java as signed, unless handled carefully by the programmer. Therefore, it is natural to make @Signed the default in the absence of a user specified annotation.




Fig. 3. The type qualifier hierarchy of our type system. Qualifiers in gray are used internally, and should not be written.

This means that the only way to introduce `@Unsigned` into a program is to annotate it explicitly. This is only problematic when a developer does not know the full specification of the program. Some of these rules lead to unnecessary imprecisions in the Signedness Type System. Later, we will identify such situations and patches to the system to fix them.

The Signedness Type System is only useful if it can identify signedness errors at compile time. To accomplish this analysis, we define several type rules. These type rules allow the Signedness Type System to identify erroneous code and issue a warning if a program might perform a computation using operands of incorrect or mixed signedness. These type rules are as follows:

- Unsigned values may not be used with sensitive operators.
- With the exception of shifts, no operator may operate on a mix of signed and unsigned values.
- Logical right shifts may only be applied to unsigned values.
- Arithmetic right shifts may only be applied to signed values.

We have already motivated the need for the first two rules, but it may come as a surprise that shifts get so much special treatment. As it turns out, the right operand of shifts is bounded within the interval $[0, 64]$. As a result, it is well within the signed positive range of even the smallest integral type, byte. Therefore, we are only interested in the left operand of shifts when applying type rules. Furthermore, since right shifts introduce higher order bits, they are dependent on the signedness of their left operand. This means that there are actually two implementations of right shift, one for unsigned integers which only introduces zeroes, and one for signed integers which depends on the sign of the integer. Our type system mandates that the correct right shift implementation is used. We will see later that this rule can be loosened slightly in specific situations.

```

// The Signedness Checker issues 3 warnings because
// unsigned right shift changes its argument's sign
// by filling in zeroes in the most significant bits.
// This doesn't matter in the below code snippet,
// because all of the introduced bits are masked off.

// c : @Signed int
sb.data[i++] = (byte) ((c >>> 8) & 0xff);
sb.data[i++] = (byte) ((c >>> 16) & 0xff);
sb.data[i++] = (byte) ((c >>> 24) & 0xff);

```

Fig. 4. An example of code for which the Signedness Checker will issue a false alarm due to masked shifts, from our case study of `jake2`.

3 PRECISION IMPROVEMENTS

As stated, there are a variety of subtle, yet important, imprecisions in the Signedness Type System. These imprecisions were discovered in the course of an evaluation of the type system on a case study, which is described in section 5. This section describes such imprecisions and improvements to the type system to solve them. All but the last of these improvements have been incorporated in an implementation of the Signedness Type System.

3.1 Masked Shifts

We recall that the type rule for right shifts is that the signedness of the left operand must match the signedness of the shift. However, when the result of the shift is masked by certain values, the introduced bits are discarded and the signedness of the shift becomes irrelevant. See Figure 4 for an example of this situation. To solve this problem, we check if every right shift is nested in a mask. If it is, we then compare the shift value and mask literal to see if the bits introduced by the shift are masked away. If they are, then we issue no error. Otherwise, we refer back to the previous behavior.

3.2 Casted Shifts

We recall again that the type rule for right shifts is that the signedness of the left operand must match the signedness of the shift. However, when the result of the shift is casted down to certain types, it could be the case that the bits introduced by the shift are discarded by the cast, rendering the signedness of the shift irrelevant to the result of the cast. See Figure 5 for an example of this situation. To solve this problem, we check if every right shift is nested in a cast. We then compare the shift value and the type of the shifted value to the type of the cast target to see if the introduced bits are discarded or not. If they are then we do not issue an error. If not, then we defer to the previous behavior.

3.3 Shift Propagation

We recall that the type of most binary operators is the least upper bound of the types of its operands. However, in the case of shifts, this is imprecise. The reason for this is that shifts are applied to the left operand, using the right operand to dictate how the shift is applied. Additionally, the bounds set by Java on the allowable values for the right operand


```

// The Signedness Checker issues a warning for this shift expression
// because an unsigned right shift is applied to a signed operand.
// However, the result of the shift is casted from an int to a byte,
// so only the 8 lower order bits matter. Since val is only shifted
// by 8 bits, it is impossible for the bits introduced by the shift
// to enter the 8 lower order bits of the result, and so the
// signedness of the shift is irrelevant.

// val : @Signed int
... (byte) (val >>> 8) ...

```

Fig. 5. An example of code for which the Signedness Checker will issue a false alarm due to casted shifts.

keep its true value well within the signed positive range. This means that only the signedness of the left operand of a shift expression is relevant to its signedness. Therefore, the type of a shift expression is the type of its left operand. See Figure 6 for an example of false errors issue by the Signedness Type System without this change.

3.4 Fully Compatible Integers

We define a “fully compatible integer” to be one which maintains the same abstract value when interpreted as signed or unsigned. Functionally, these are integers which are in the signed positive range for a given bit width. Another, more useful way of expressing this quality is that these are integers whose most significant bit is zero. Such integers can be safely interpreted as signed and unsigned, and indeed it is common to use such integers with less discretion than other integers. As such, not accounting for this fact is a huge source of imprecision in the Signedness Type System. If a variable is known to always be in the signed positive range, then it can safely be used as signed or unsigned. In its current state, the Signedness Type System is not powerful enough to infer this property, and therefore considers every integer to be either @Signed or @Unsigned, unless it can infer it to be @Constant. See Figure 7 for an example of a false error produced when a fully compatible integer is used as both @Signed and @Unsigned. To solve this problem, we plan to make use of another type system [2] for representing value ranges, in order to gather information about each variable for new introduction rules. We then refine variables which are always within their signed positive range as @Constant. More specifically, the proposed introduction rules follow:

- byte b is @Constant if it has type @IntVal{vs} and each v in vs satisfies $0 \leq v < 2^7$.
- char c is @Constant if it has type @IntVal{vs} and each v in vs satisfies $0 \leq v < 2^7$.
- short s is @Constant if it has type @IntVal{vs} and each v in vs satisfies $0 \leq v < 2^{15}$.
- int i is @Constant if it has type @IntVal{vs} and each v in vs satisfies $0 \leq v < 2^{31}$.
- long l is @Constant if it has type @IntVal{vs} and each v in vs satisfies $0 \leq v < 2^{63}$.
- byte b is @Constant if it has type @IntRange{a, b} and a and b satisfy $0 \leq a \wedge b < 2^7$.
- char c is @Constant if it has type @IntRange{a, b} and a and b satisfy $0 \leq a \wedge b < 2^7$.
- short s is @Constant if it has type @IntRange{a, b} and a and b satisfy $0 \leq a \wedge b < 2^{15}$.
- int i is @Constant if it has type @IntRange{a, b} and a and b satisfy $0 \leq a \wedge b < 2^{31}$.
- long l is @Constant if it has type @IntRange{a, b} and a and b satisfy $0 \leq a \wedge b < 2^{63}$.

```

// The Signedness Checker issues a warning because the expression
// (vis[cluster>>3] & 0xFF) has type @Unsigned, but the expression
// (1 << (cluster & 7)) has type @Signed when really it could be
// interpreted as signed or unsigned, and should have type @Constant.

// vis : @Unsigned byte[], cluster : @Signed int
(vis[cluster>>3] & 0xFF) & (1 << (cluster & 7))

```

Fig. 6. An example of code for which the Signedness Checker will issue a false alarm due to shift propagation, from our case study of `jake2`.

4 BACKGROUND AND IMPLEMENTATION

Now that we have defined a type system, the Signedness Type System, to identify signedness errors, we wish to implement a tool, the Signedness Checker, for applying this type system to Java programs. Our goal is to build a verification tool that guarantees that software is free of bugs related to unsigned integers. To achieve this, our approach must be sound: if it issues no warnings, then the program must be free of bugs. Any sound analysis sometimes issues false positives — warnings about code that will not go wrong at run time. This occurs when the Signedness Checker cannot prove that the code is correct. We have identified several sources of imprecision in the Signedness Type System, and therefore the Signedness Checker. Most of these imprecisions have been fixed in the latest version of the tool.

We built our Signedness Checker implementation upon the Checker Framework [4, 9], which enables the construction of pluggable type systems for Java. The Signedness Checker is distributed as part of the Checker Framework (<http://CheckerFramework.org/>). See Figure 8 for a depiction of the workflow of this system.

Our pluggable type system implementation consists of roughly 324 statements. One of its largest components is a run-time library that provides JDK 8 unsignedness functionality to earlier versions of Java. The type-checker consists of the following parts:

- A series of type qualifiers implemented as Java 8 type annotations [6].
- An Annotated Type Factory where type introduction rules are defined procedurally.
- An AST Visitor which enforces type rules during a traversal of the abstract syntax tree.

```

// The Signedness Checker issues a warning because
// CL_fx.ParticleEffect receives an unsigned third argument.
// However, the variable color is always within the range
// [0, 255] inclusive, so it is well within the int signed
// positive range.

// color : @Unsigned int
// splash_color = { 0x00, 0xe0, 0xb0, 0x50, 0xd0, 0xe0, 0xe8 }
if (r > 6)
    color = 0x00;
else
    color = splash_color[r];

// CL_fx.ParticleEffect(float, float, @Signed int, @Signed cnt)
CL_fx.ParticleEffect(pos, dir, color, cnt);

```

Fig. 7. An example of code for which the Signedness Checker will issue a false alarm due for fully compatible integers.

5 RESULTS AND CONTRIBUTIONS

We evaluated our tool by running it on `jake2` [1], a Java port of the popular '90s video game, Quake II. The original C implementation of Quake II used unsigned integers. `jake2` tries to mimic Quake II's usage and consists of 133513 lines of code that can be found at (<https://github.com/mbien/jake2>).

5.1 First Approach

Our initial estimate of the codebase seemed to indicate that only a small subset of the codebase made use of unsigned integers, covering roughly 10000 lines of code. We wrote 33 annotations, referring to the limited documentation in `jake2` whenever we had questions about the developers' intent.

When running the Signedness Checker on the annotated `jake2`, we received a staggering number of false positive errors related to masked shifts and casted shifts. See Figure 4 and Figure 5 for examples of the sites of such false positive errors.

We also identified one bug (Figure 9): unsigned integers are printed by a function that does no special handling of unsigned values. Thus, they are printed as if signed, which leads to erroneous output for values outside the signed positive range.

5.2 Second Approach

After we had annotated and checked this subset of `jake2`, we decided to check our annotations against the original C implementation, Quake II. We found that the documentation in `jake2` was a drastic underestimation of the true behavior of the program. That is, we found that there were many sections of code present in `jake2` which were semantically equivalent to sections of code in Quake II, except that there was no documentation concerning variables which were




Fig. 8. Our type system is implemented in the Signedness Checker, which hooks into a modified form of javac to check the signedness property during standard compilation.

explicitly unsigned in the C implementation. We therefore decided to do another, more ambitious, case study of *jake2*.

Our plan for the second *jake2* case study was to annotate Java variables as `@Unsigned` if they had a C equivalent which had the type unsigned. This approach highlighted the unsoundness of the C type system with respect to this property. In C, no compiler errors or warnings are issued when signed and unsigned integers are mixed in a computation. Rather, signed integers are simply promoted to unsigned integers. As we’ve seen, this can lead to a variety of errors. Furthermore, due to this property of C, *Quake II* would not type check under the Signedness Type System, as many variables which are actually intended to be unsigned are not annotated as such. As a result, running the Signedness Checker on *jake2* after it has been annotated in this manner produced 1071 errors.

A brief inspection of these errors seemed to indicate that vast majority were due to the incomplete specification provided by *Quake II*. Our solution for this was to step through erroneous code and annotate variables as `@Unsigned` if we could confirm that they were used exclusively as `@Unsigned` integers. Throughout this process, we discovered that there were many errors which were not in fact due to the insufficient specification, but were false positives produced by erroneous shift propagation, as discussed earlier. See Figure 6 for an example of such false positives. After correcting this issue, we continued with our analysis of *jake2*. We eventually reached a point in the case study where it became clear that the Signedness Type System was not sufficiently precise with regard to fully compatible integers, as previously discussed. See Figure 7 for an example of a false positive caused by this imprecision. In the future, we plan to integrate the Value Checker [2] with the Signedness Checker in order to add introduction rules to capture integer ranges in our analysis.

```

// out.firstleafbrush and out.numleafbrushes are
// unsigned. Vargs::add(int) expects a signed int.

if (debugloadmap) {
    Com.DPrintf("%8x|%6i|%6i|%6i|\n",
        new Vargs()
            .add(out.contents)
            .add(out.cluster)
            .add(out.area)
            .add(out.firstleafbrush)
            .add(out.numleafbrushes));
}

```

Fig. 9. Buggy code in jake2 that our tool identified.

During the course of this expanded case study, we also identified several new bugs. All of these bugs were similar in form to the bug identified in the first case study. That is, all bugs found involved the erroneous conversion of unsigned integers to strings. See Figure 10 for an example of the new bugs found by the Signedness Checker.

5.3 Tool Evaluation

We performed the first case study in parallel with development and enhancement of the type system, spending a few hours a week over the course of several months. Almost all of our time on the case study was spent reverse-engineering the poorly documented codebase to determine the intention of its developers. Using the Signedness Checker was simple: it only requires a developer to write the `@Unsigned` annotation where they would write `unsigned int` in C, then run `javac` with an annotation processor to type-check the code. We believe a developer already familiar with the `jake2` codebase could have learned the Signedness Type System, written the 33 annotations, and discovered the bug in well under a week.

The second case study was a significantly more time-consuming endeavor, although more valuable. We spent a few hours a week over the course of a year running this case study. The vast majority of this time was spent inferring the true specification of `jake2`, in absence of an acceptable specification from Quake II. While the time commitment was much higher, it was still relatively simple to use the Signedness Checker. However, it did become necessary to annotate all final static integral fields as `@Constant` in addition to the usual `@Unsigned` annotations as the Signedness Checker does not infer field types. In total we applied 528 `@Unsigned` annotations and 7352 `@Constant` annotations. Note that we were able to apply all 7352 `@Constant` annotations with a single find and replace pass, whereas the 528 `@Unsigned` annotations took months of manual inspection. We were able to verify a full module, the `sys` module, using the Signedness Checker, and we were also able to identify at least 5 bugs using the Signedness Checker. Furthermore, we were able to discover the Signedness Checker’s erroneous shift propagation and lack of precision with respect to fully compatible integers through this case study.

This experience has allowed us to conclude that the Signedness Checker is probably more useful for well documented code, or during development of a new application. If a code base is poorly specified with respect to signedness,

```

// out is an unsigned array.
// Vargs::add(int[]) expects a signed array.

for (int i = 0; i < count; i++) {
    out[i] = SignednessUtil.getUnsignedShort(bb);
    if (debugloadmap) {
        Com.DPrintf("|%6i|%6i|\n",
            new Vargs()
                .add(i)
                .add(out[i]));
    }
}

```

Fig. 10. More buggy code in jake2 that our tool identified.

the Signedness Checker cannot infer intent, as it is the responsibility of the developer to introduce @Unsigned annotations.

6 RELATED WORK

SmartFuzz [7] generates tests that are intended to expose integer bugs, including signed/unsigned conversion. For each program execution, it generates and solves a set of constraints that assigns each integer to a lattice containing Top, Signed, Unsigned, and Bottom. SmartFuzz incorrectly under-constrains some operations such as x86’s IMUL. It uses Valgrind to partition integers into union-find sets and garbage-collects them, an implementation strategy pioneered by DynComp [5]. They report finding 5 likely distinct bugs in C programs. As with any testing tool, SmartFuzz gives no guarantee about future executions. The same authors’ CatchConv tool [8] also does run-time type inference, focusing on integer conversion errors and utilizing symbolic execution. Our approach is more lightweight and precise and gives a guarantee, but it requires programmers to write a few annotations in their code.

7 FUTURE WORK

Currently, we have proposed a fix to the imprecision concerning fully compatible integers in the Signedness Type System. To reiterate, we plan to refine the type of such integers to @Constant by using the Value Type System [2] to allow us to infer when an integer is fully compatible. We also plan on allowing the Signedness Type System to capture the signedness of boxed integral types, special Java wrappers for integral primitives.

8 CONCLUSIONS

We have developed and implemented a type system, the Signedness Type System, that captures usage of signed and unsigned integers in Java programs. This type system enables developers to detect errors regarding unsigned integers at compile time and guarantees that such errors cannot occur at run time. In a case study, an implementation of the type system proved easy to use and helped uncover 5 previously unknown bugs.

REFERENCES

- [1] 2009. jake2. <https://github.com/mbien/jake2>. (Dec. 2009).
- [2] 2017. *Constant Value Checker*.
- [3] Joseph D. Darcy. 2012. Unsigned Integer Arithmetic API now in JDK 8. https://blogs.oracle.com/darcy/entry/unsigned_api. (Jan. 2012).
- [4] Werner Dietl, Stephanie Dietzel, Michael D. Ernst, Kivanç Muşlu, and Todd Schiller. Building and using pluggable type-checkers. 681–690.
- [5] Philip J. Guo, Jeff H. Perkins, Stephen McCamant, and Michael D. Ernst. Dynamic inference of abstract types. 255–265.
- [6] JSR 308 Expert Group. 2014. Annotations on Java types. http://download.oracle.com/otndocs/jcp/annotations-2014_01_08-pfd-spec/. (Jan. 8, 2014). Proposed Final Draft.
- [7] David Molnar, Xue Cong Li, and David A. Wagner. Dynamic test generation to find integer bugs in x86 binary Linux programs. 67–82.
- [8] David Alexander Molnar and David Wagner. 2007. *Catchconv: Symbolic execution and run-time type inference for integer conversion errors*. Technical Report UCB/EECS-2007-23. EECS, UC Berkeley, Berkeley, CA, USA.
- [9] Matthew M. Papi, Mahmood Ali, Telmo Luis Correa Jr., Jeff H. Perkins, and Michael D. Ernst. Practical pluggable types for Java. 201–212.